

Input Layer
PDF Document — User Upload

PDF Loader (PyMuPDF)

Text — Table — Figure Extractors

Semantic Chunker
512 tokens, 50 overlap

Embedding Model
all-MiniLM-L6-v2

ChromaDB Vector Store

BM25 Keyword Index

User Query

Query Parser + Intent Detection

Hybrid Search
Vector + Keyword

RRF Fusion

Cross-Encoder Reranker

Gemini 2.0 Flash

Rule-Based Fallback

Answer + Citations

Streamlit UI

Evaluation Suite